


Intro to Reinforcement Learning


Markovich Anna

Who is this **agent** who uses **reinforcement** to beat the **environment**?

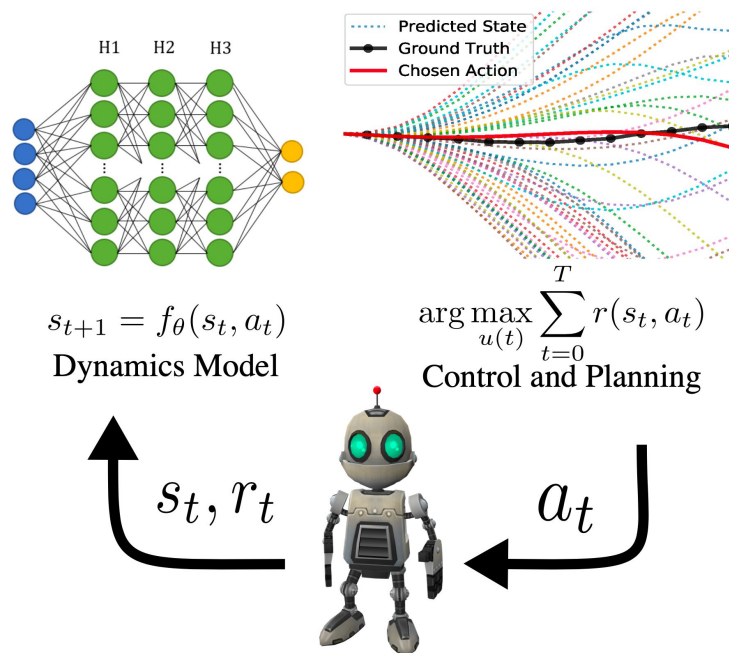


Environment | State | Action | Reward

Game: tetris



Neural network



RecSys



YouTube

Problem: how to learn optimal behaviour?

1. Exploration
 2. Optimization
 3. Generalization
- handle delayed consequences

Let's compare:

	Reinforcement learning	Supervised learning
Exploration	trial and error	<u>None</u> (known ground truth)
Optimization	sparse reward function	differentiable loss function
Generalization	update decision policy	avoid overfitting
Delayed consequences	reward may come afterwards	<u>None</u>

Formalism: Markov Decision Process

Markov assumption:

$$P(s_{t+1} | s_t, s_{t-1}, s_{t-2} \dots s_1) = P(s_{t+1} | s_t)$$

(next state depends only on the current state)

Formalism: Markov Decision Process

Markov chain / Markov process – a **stochastic model** describing a **sequence** of possible events which satisfies Markov assumption.

(S, P) :

S – state space

$$s \in S$$


P – dynamics (conditional distribution)

$$P(s', s) = P(s_{t+1} = s' | s_t = s)$$

Formalism: Markov Decision Process

MDP = Markov chain + action + reward + policy

$$P(s', s, a) = P(s_{t+1} = s' | s_t = s, a_t = a) \quad \pi(a|s) = P(a_t = a | s_t = s)$$



$$R(s, a) = \mathbb{E}[r | s_t = s, a_t = a]$$

Total reward / Return

total rewards for session

$$G_t = \sum_{s=t}^T R_s$$



with discount

$$G_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^T R_{t+T}$$

$$\gamma \in (0, 1)$$

Goal

find policy :

$$\mathbb{E}_{\pi}[G_t] \rightarrow \max_{\pi}$$

Cross-entropy method

0. initialize policy

loop:

1. sample N sessions
2. choose $M < N$ **elite** sessions
3. update policy to favor pairs (a, s) from **elite**

Cross-entropy method

Tabular : finite state & action spaces



PyTorch
Tabular



policy matrix: $\pi(a|s) = \Pi_{a,s} \in \mathbb{R}^{|A| \times |S|}$

update rule:
$$\pi(a_p|s_k) = \frac{\sum_{(a,s) \in elite} [s_k = s][a_p = a]}{\sum_{(a,s) \in elite} [s_k = s]}$$

Cross-entropy method

Infinite/large state space

Maximum Likelihood
Estimation


Methods of Economic
Investigation
Lecture 17

loop:

1. sample N sessions
2. choose M elite sessions
3. $\pi = \underset{\pi}{argmax} \sum_{(a_i, s_i) \in elite} \log(\pi(a_i | s_i))$

Cross-entropy method

Infinite/large state space



Using NNs:

0. `net = MLPClassifier(...)`

loop:

1. sample N sessions
2. choose M elite sessions
3. `net.fit(elite_states, elite_actions)*`

$$*W_k = W_{k-1} + \eta_k \nabla \left(\sum_{(a_i, s_i) \in \text{elite}} \log(\pi_{W_{k-1}}(a_i | s_i)) \right)$$

Cross-entropy method

Continuous state & action space

model: $\pi(a|s) \sim \mathcal{N}(\mu(s), \sigma^2)$

μ - NN output

σ - parameter / other NN output

Using NNs:

0. **net = MLPRegressor(...)**

loop:

1. sample N sessions
2. choose M elite sessions
3. **net.fit(elite__states, elite__actions)***

$$*W_k = W_{k-1} + \eta_k \nabla \left(\sum_{(a_i, s_i) \in \text{elite}} \log(\pi_{W_{k-1}}(a_i | s_i)) \right)$$

Intrigue: better method for optimal policy search

(not today)

Remember ?

$$G_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^T R_{t+T}$$

$$\mathbb{E}_\pi[G_t] \rightarrow \max_\pi$$

V-function and Q-function

State-value function

$$V_{\pi}(s) = \mathbb{E}_{\pi}[G_t | s_t = s]$$

expected return
conditional on state

Action-value function

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | s_t = s, a_t = a]$$

expected return
conditional on state and action


Complex mathematical calculations

$$V_{\pi}(s) = \sum_a \pi(a|s) \sum_{r,s'} p(r, s' | s, a) [r + \gamma V_{\pi}(s')] = \sum_a \pi(a|s) Q_{\pi}(s, a)$$

$$Q_{\pi}(s, a) = \sum_{r,s'} p(r, s' | s, a) [r + \gamma V_{\pi}(s')]$$

Policy iteration, Value iteration

Policy evaluation



Policy improvement



1. deterministic policy :

$$\pi : S \rightarrow A$$

2. $\pi(s) \leftarrow \operatorname{argmax}_a Q_{\pi}(s, a)$

Sources

https://github.com/yandexdataschool/Practical_RL

<http://web.stanford.edu/class/cs234/>