

# Tensor Programming V

Denis Sapozhnikov

HSE AMI 202


January 15, 2024

- ① Motivation
- ② Intuition
- ③ Why SP is bad?
- ④ The  $\mu P$  formulas
- ⑤ Experiments
- ⑥ Discussion

# Large Models problems

- Speed up of hyperparameter tuning (HP)
- Cost of HP
- Quality of HP

# Solution



# Central Limit Theorem

Kind reminder about the CLT formula. If  $x_1, x_2, \dots, x_n$  "look like" sample of random independent variable  $X$ , then

$$\frac{1}{\sqrt{n}} \sum_{i=1}^n (x_i - \mathbb{E}[X]) \rightarrow \mathcal{N}(0, \sigma(X))$$

We will call  $\frac{1}{\sqrt{n}}$  the right scaling factor  $c_n$ , because with its value the formula yields non-trivial distribution.

# Minimization Task

Now define some minimization problem  $F_n(c) \rightarrow \min$  like following:

$$F_n(c) = \mathbb{E}_{x_1, \dots, x_n} f(c(x_1 + x_2 + \dots + x_n)),$$

where  $x_1, x_2, \dots, x_n$  are hidden variables and  $f$  is a bounded function.

Here we can define the right scaling factor  $c_n = \frac{\alpha}{\sqrt{n}}$ , because we would minimize something non-trivial in infinite-width case:

$$\lim_{n \rightarrow \infty} F_n(c_n) \rightarrow f(\mathcal{N}(0, \alpha^2)) =: G_n(\alpha) \quad (1)$$

# Ok, and?

- The equation (1) means that for sufficiently large  $n$ , the optimal  $\alpha_n^* = \arg \min G_n(\alpha)$  should be close to  $\alpha_N^*$  for any  $N > n$ .
- So, applying this to the ideas of machine learning, we can select the scaling factor  $c_n$  (learning rate, width, etc.) so that for all larger models the model quality is optimal without HP.

# MLP with one hidden layer

Let's define a standard MLP with one hidden layer  $v$  of width  $n$ , input layer  $u$  ( $u, v \in \mathbb{R}^n$ ), and 0 biases for one scalar sample  $x$ :

$$f(x) = v^T u x$$

with a standard parametrization (SP):  $v_i \sim \mathcal{N}(0, \frac{1}{n})$  and  $u_i \sim \mathcal{N}(0, 1)$  and learning rate  $\eta$ .



# Why SP is bad

After the first step of SGD, the updated weights will look like  $v \leftarrow v + \theta u, u \leftarrow u + \theta v$ . From now on the function  $f(x)$  is following:

$$f(x) = (v + \theta u)^T (u + \theta v)x = (v^T u + \theta u^T u + \theta v^T v + \theta^2 u^T v)x$$

As you can mention,  $u^T u \in \Theta(n)$  which is blown up with the width of the network. Hence, the model's weights will explode in an infinite-width case.

# A brand new parametrization

Let's fix our parametrization with a new one ( $\mu P$ ):

- $v_i \sim \mathcal{N}(0, 1), u_i \sim \mathcal{N}(0, \frac{1}{n^2})$
- $\eta_v = \frac{1}{n}\eta, \eta_u = n\eta$

Then after updating the weights formula will look like this:

$$f(x) = (v^T u + \theta n^{-1} u u^T + \theta n v^T v + \theta^2 u^T v) x$$

Why is it better?

# Proof

- $n^{-1} \mathbb{E}[u^T u] = n^{-1} \cdot n \cdot \mathbb{E}u_0^2 = 1 \cdot 1 = 1$
- $n \mathbb{E}[v^T v] = n^2 \mathbb{E}v_0^2 = n^2 \cdot \frac{1}{n^2} = 1$
- $\mathbb{E}[u^T v] = n \cdot \mathbb{E}[u_0 v_0] \leq [\text{Cauchy-Schwarz}] \leq n \cdot \sqrt{\text{Var}(v_0)} \sqrt{\text{Var}(u_0)} = n \cdot 1 \cdot \frac{1}{n} = 1$


So the weights in  $\Theta(1)$ , therefore there are no vanishing or exploding of model weights.

# How to fix any layer?

	Input weights & all biases	Output weights	Hidden weights
Init. Var.	$1/\text{fan\_in}$	$1/\text{fan\_in}^2 (1/\text{fan\_in})$	$1/\text{fan\_in}$
SGD LR	$\text{fan\_out} (1)$	$1/\text{fan\_in} (1)$	1
Adam LR	1	$1/\text{fan\_in} (1)$	$1/\text{fan\_in} (1)$


But many details about other parameters and special layers (e.g. Transformers) are in the appendix.

# MLP



- $\mu P$  works on "trivial" networks
- Wider is better for any training step
- Unlike SP, optimum has no shift with a rising width

## GPT-3



- The proxy model is 168 times smaller than the target model by reducing the width by a factor of 16
- The proxy-model was trained only on 4 or 16 billion tokens while the target used the 300 billion
- Tuning cost only 7% of total pretraining cost
- The model exceeded the results of the original work and is comparable to models 2 times larger

# Results

- Better performance:  $\mu P$  outperforms SP
- Theory is working on practice with all tested families of models
- Unlike previous works, there is quite a bit family of "broken" layers in scaling rules
- Semi-automatic framework by authors could reduce your pain
- Now more researchers can afford to experiment with large models, and comparing the quality of models will be much easier and more convenient with the same approach to finding hyperparameters

# Criticism

- Theory and practice still require a model of a “sufficiently large” size, while there is no understanding of what size is enough
- Based on the plots, the optimal SP models can learn faster than  $\mu P$
- The optimal HP still shifts slightly for smaller models
- Initialization does not transfer well across depth, and depth transfer generally still does not work for post-layernorm Transformers